

Research Notes: A New Surrogate Closure Technique for the M_N System

William Porteous

Oak Ridge National Laboratory: CSM Department

December 30, 2020

Contents

1	Introduction	2
1.1	The Closure Framework for Kinetic Equations	2
1.2	Entropy-Based Closures: The M_N System	2
1.3	The Lagrange-Dual to Entropy-Minimization	3
1.4	The Computational Problem: An Explicit Example	4
2	Proposed Method	4
2.1	Metric for Approximation Accuracy	5
2.2	Dimension-Reduction: From Unbounded to Bounded Regions	5
2.2.1	Scaling Relations	5
2.2.2	Guaranteeing Convexity	6
3	Convexity of Scaled-Approximations	6
3.1	Function-Based Approximation: Method A	6
3.2	Gradient-Based Approximation: Method B	6
3.3	Convexity of Method A	7
3.4	Convexity Question of Method B	8

These notes motivate and introduce my research at ORNL-CSM, and do not divulge our numerical results. For brevity, examples are built on a 1-dimensional velocity domain. We examine a velocity-discretized modification of a kinetic equation, the M_N equations, which we intend to solve numerically with a Runge-Kutta scheme. We are interested in improving the efficiency of this solution scheme by introducing a new technique to calculate the flux-update, a key step in the scheme, and are especially concerned with the efficiency of this method as N (a positive integer) increases. Our technique will employ the same Lagrange-duality framework of M_N to calculate the flux-update, but will supplant the $N + 1$ dimensional optimization problem of this framework with an N -dimensional approximation problem.

1 Introduction

1.1 The Closure Framework for Kinetic Equations

Consider a generic kinetic equation for a system of charge neutral particles which interact with each other through collision, and perhaps with a background medium. This kinetic equation evolves a kinetic density function $f: [0, \infty) \times X \times V \longrightarrow [0, \infty)$ according to:

$$\partial_t f(t, x, v) + v \cdot \nabla_x f(t, x, v) = \mathcal{C}(f(t, x, \cdot))(v) \quad (1.1.1)$$

where \mathcal{C} , appropriately called the "collision operator", introduces all interaction effects. In order to be well-posed, this equation requires specific initial and boundary conditions. With the intent of discretizing (1.1.1) in the velocity variable, we consider an orthonormal basis $\{\varphi_n\}_{n \in \mathbb{N}}$ for $L^2(V)$ and make the definitions

$$\begin{aligned} \mathbf{m}_N(v) &= (\varphi_0(v), \varphi_1(v), \dots, \varphi_{N-1}(v)) \\ \mathbf{u}: [0, \infty) \times X &\longrightarrow \mathbb{R}^{N+1} \end{aligned}$$

$$\mathbf{u}^i(t, x) := \int_V f(t, v, x) \mathbf{m}_N^i(v) dv \text{ for } 0 \leq i \leq N$$

The function \mathbf{u} is a vector containing the first $N + 1$ moments of f according to the moment basis functions contained in \mathbf{m} . The moments \mathbf{m} induce a non-linear Galerkin discretization of (1.1.1) along the velocity domain as

$$\partial_t \langle \mathbf{m} F_{\mathbf{u}} \rangle + \nabla_x \langle v \cdot \mathbf{m} F_{\mathbf{u}} \rangle = \langle \mathbf{m} \mathcal{C}(F_{\mathbf{u}}) \rangle \quad (1.1.2)$$

where $F_{\mathbf{u}(t,x)}$ is some approximation to f with the same velocity moments as f at each point in space and time. That is to say, we demand that $F_{\mathbf{u}(t,x)}$ is some ansatz for f which satisfies the constraint:

$$\int_V F_{\mathbf{u}(t,x)}(t, x) \mathbf{m}(v) dv = \mathbf{u}(t, x)$$

Of course, we want this ansatz F to be physically "reasonable", satisfying properties of our original kinetic equation (1.1.1) like conservation of mass and the H-Theorem.

1.2 Entropy-Based Closures: The M_N System

One such form for F which preserves physically-desireable properties of (1.1.1) is the entropy-based closure, where we assume F is a nonnegative function with minimal total entropy along the velocity domain with the moments \mathbf{u} . Let's make this notion precise, and for the sake of notation, henceforth write $\langle f \rangle$ to mean $\int_V f dv$.

Consider a convex, C^2 function, $\eta: \mathbb{R}_+ \longrightarrow \mathbb{R}$, called the kinetic entropy density.

Definition (Entropy-Based Closure). *In the case of (1.1.2), we demand that $F_{\mathbf{u}(x,t)}$ is the solution to the following minimization problem:*

$$\begin{aligned} &\text{minimize } \langle \eta(g) \rangle, \quad g \in L^1(V) \\ &\text{subject to } \langle \mathbf{m} g \rangle = \mathbf{u}(x, t) \end{aligned}$$

Where we let take η to be the Boltzmann entropy, $\eta(r) = r \log(r) - r$, and restrict to $N + 1$ moments, i.e. $\mathbf{u}(x, t) = \langle f \cdot \mathbf{m}_N \rangle$, this is called the M_N closure framework, and the corresponding expression for (1.1.2) is called the M_N system.

Definition (M_N System). In (1.4.2), define the function $F(\mathbf{u}_N)$ as the extremal function of the following minimization (entropy-maximization) problem

$$\begin{aligned} & \text{minimize } \langle \eta(g) \rangle, \quad g \in L^1(V) \\ & \text{subject to } \langle \mathbf{m}_N g \rangle = \mathbf{u}_N \end{aligned} \quad (1.2.1)$$

where $\langle \cdot \rangle$ denotes integration with respect to the velocity domain V , and $\eta(r) = r \log(r) - r$.

Definition (The Realizable Set). For the M_N framework, we denote the set of all possible moment vectors according to V and \mathbf{m} by

$$\mathcal{R} := \{\mathbf{u} \in \mathbb{R}^{N+1} : \mathbf{u} = \langle f \mathbf{m}_N \rangle, \quad f \in L^1(V)\}$$

1.3 The Lagrange-Dual to Entropy-Minimization

The problem may have no solutions, or fail uniqueness, in the case that V is unbounded. There are explicit examples which saturate each of these. To continue, we assume that V is a convex, compact subset of \mathbb{R} . For our purposes, we will often consider $V = [-1, 1]$, and in that case, we will take that \mathbf{m}_N is the vector of the first $N + 1$ Legendre polynomials (including the 0th order identity map).

Definition (Minimum Entropy Function).

$$h : \mathcal{R} \longrightarrow \mathbb{R}$$

$$h(\mathbf{u}) = \min \langle \eta(g) \rangle \text{ subject to } g \in L^1(V), \quad \langle \mathbf{m} g \rangle = \mathbf{u} \quad (1.3.1)$$

The above minimization problem of the M_N system has the following unconstrained dual problem (writing η_* for the legendre-transform of η):

$$\text{minimize } \langle \eta_*(\alpha^t \mathbf{m}) \rangle - \alpha^t \cdot \mathbf{u}, \quad \alpha \in \mathbb{R}^{N+1} \quad (1.3.2)$$

where $h_*(\alpha) = \langle \eta_*(\alpha^t \mathbf{m}) \rangle$.

Definition (Dual Solution $\hat{\alpha}(\mathbf{u})$). For an admissible moment vector u

$$\hat{\alpha}(\mathbf{u}) := \operatorname{argmin} \langle \eta_*(\alpha^t \mathbf{m}) \rangle - \alpha^t \cdot \mathbf{u}, \quad \alpha \in \mathbb{R}^{N+1}$$

Fact (Primal Solution). By necessary conditions for duality, the solution $\hat{\alpha}(\mathbf{u})$ to the dual problem obtains the primal solution to (1.3). Explicitly, the function $G_\alpha(v)$ defined as

$$G_\alpha(v) = \eta'_*(\alpha^t \mathbf{m}) = \exp(\alpha^t \mathbf{m}) \quad (1.3.3)$$

is the solution to (1.3)

Fact (Gradient is Dual Optimum). First order conditions for duality show that

$$\nabla h(\mathbf{u}) = \hat{\alpha}(\mathbf{u}) \quad (1.3.4)$$

From the above fact, it immediately follows that $\hat{\alpha}$ is therefore an invertible function, let its inverse be given by $\hat{u}(\alpha)$ for admissible vectors $\alpha \in \mathbb{R}^{N+1}$.

1.4 The Computational Problem: An Explicit Example

Consider the case $N = 1$ and the following setup:

$$V = [-1, 1], \quad \mathbf{m} = \begin{pmatrix} 1 \\ v \end{pmatrix}$$

Note that the functions in \mathbf{m} are orthogonal, but not normalized in $L^2(V)$. We define a kinetic distribution function $f = f(x, \mu, t)$ for $x \in \mathbb{R}, \mu \in V = [-1, 1], t \geq 0$ which evolves according to the following **linear kinetic equation**

$$\partial_t f + \mu \partial_x f + \sigma_a f = \sigma_s \left(\frac{1}{2} \langle f \rangle - f \right) \quad (1.4.1)$$

We write out what this means for our Galerkin discretization (1.1.2):

$$\partial_t \mathbf{u}_N + \partial_x \langle \mu \mathbf{m}_N f \rangle + \sigma_a \mathbf{u}_N = -\sigma_s \mathbf{Q} \mathbf{u}_N$$

This is not "closed" in the first $N + 1$ moments of f , since the second term on the lefthand-side contains higher order moments of f . We close this as an evolution of \mathbf{u}_N with our M_N technique, giving:

$$\partial_t \mathbf{u}_N + \partial_x \langle \mu \mathbf{m}_N F(\mathbf{u}_N) \rangle + \sigma_a \mathbf{u}_N = -\sigma_s \mathbf{Q} \mathbf{u}_N \quad (1.4.2)$$

To numerically solve a discretization of the PDE (1.4.2) via a Runge-Kutta scheme, we must compute $G(\mathbf{u}_N)$ via $G_\alpha(v)$, which is defined by (1.3.2). Each update to the numerical \mathbf{u}_N , requires the computation of $\hat{\alpha}(\mathbf{u})$. This can be done accurately for each value of \mathbf{u}_N , for example, via a gradient-descent optimization algorithm. This method has the advantage that, with high accuracy for values of α , we are assured that $\hat{\alpha}(\mathbf{u})$ are inherited from the convex function h , ensuring our solution enjoys a variety of desirable properties. However, this is computationally slow – the optimization routine becomes much slower as the number of moments (N) increases, and similarly, if the velocity and position domains increase in dimension.

2 Proposed Method

We can approximate the entropy function h with a function \tilde{h} , which is convex and C^2 , but which also admits a closed-form, forward-evaluation expression for its gradient $\hat{\alpha}$.

1. We approximate the entropy function $h: \mathcal{R} \rightarrow \mathbb{R}$ with \tilde{h} using
 - (a) A fully-connected artificial Neural-Network
 - (b) Rational-Cubic Spline
2. Compute $\nabla \tilde{h}(\mathbf{u})$ and define $\tilde{\alpha} := \nabla \tilde{h}(\mathbf{u})$

In principal, this can be done on the full 2 and 3-dimensional realizable sets for the M_1 equation and M_2 equations. However, we were able to derive scaling relations which enabled us to reduce the problem to approximating h on the interior of 1-dimensional segment, and on the interior of a parabola.

2.1 Metric for Approximation Accuracy

The term which we desire to compute more efficiently in (1.4.1) is the so-called "flux-update", i.e.

$$F: \mathcal{R} \longrightarrow \mathbb{R}^2$$

$$F(\mathbf{u}) := \langle \mu \mathbf{m}_N G_{\tilde{\alpha}(\mathbf{u})} \rangle$$

Fact. *It follows from calculus that*

$$\|F(\mathbf{u}) - F(\mathbf{v})\| \leq C \|\mathbf{u} - \mathbf{v}\|$$

For an approximate closure $\tilde{h}, \tilde{\alpha}$, \mathbf{u} and v are replaced in this estimate $\hat{u}(\tilde{\alpha}(u))$ and by $\hat{u}(\tilde{\alpha}(v))$ - the moment-reconstruction induced by our approximation to α . Since the relationship between $\tilde{\alpha}$ and $\hat{u}(\tilde{\alpha})$ is defined by the integral of an exponential, we consider the $\hat{u}(\tilde{\alpha}(u))$ to be the best metric of our accuracy.

For accuracy, we intend to obtain an approximation which minimizes the value

$$\|\hat{u}(\tilde{\alpha}(u)) - u\|_{L^2(\mathcal{R})}$$

2.2 Dimension-Reduction: From Unbounded to Bounded Regions

2.2.1 Scaling Relations

From the definitions in Section 1, it is straightforward to show that

Fact (Alpha Scaling).

$$\begin{aligned} \forall \lambda > 0 \\ \hat{\alpha}(\lambda \mathbf{u}) = \alpha(\mathbf{u}) + \begin{pmatrix} \ln(\lambda) \\ 0 \end{pmatrix} \end{aligned} \quad (2.2.1)$$

from which it follows that

Fact (Entropy Scaling).

$$\begin{aligned} \forall \lambda > 0 \\ h(\lambda \mathbf{u}) = \lambda(h(\mathbf{u}) + u_0 \ln(\lambda)) \end{aligned} \quad (2.2.2)$$

For the case of the M_1 equation, where the moment vectors are simply $\mathbf{u} = (u_0, u_1)$, this gives the following:

$$\begin{aligned} \textbf{Fact } (u_0 \text{ Scaling}). \text{ For } v \in \mathcal{R}, \text{ choose } \lambda = v_0 \text{ and } \mathbf{v} = v_0 \begin{pmatrix} 1 \\ \frac{v_1}{v_0} \end{pmatrix} \\ h(\mathbf{v}) = v_0(h(\frac{\mathbf{v}}{v_0}) + \ln(v_0)) \end{aligned} \quad (2.2.3)$$

To approximate h , we can approximate the behavior of h on the set $\mathcal{D} = \{v \in \mathcal{R}: v_0 \equiv 1\}$, and extend the network to \mathcal{R} using this scaling relation. That is

Definition. *We write \tilde{h}_s for a function*

$$\tilde{h}_s: \mathcal{D} \longrightarrow \mathbb{R}$$

which approximates the entropy function h on the set \mathcal{D}

$$\tilde{h}: \mathcal{R} \longrightarrow \mathbb{R}$$

$$\tilde{h}(\mathbf{u}) := u_0 \cdot \tilde{h}_s((1, u_1/u_0)) + u_0 \cdot \ln(u_0)$$

2.2.2 Guaranteeing Convexity

It is not immediately obvious (to me) that an approximation \tilde{h}_s which is convex will induce a convex function \tilde{h} . This, however, turns out to be the case; we prove this by checking the positive-definiteness of the hessian matrix for \tilde{h} .

3 Convexity of Scaled-Approximations

In both methods here, a neural network approximation to the entropy function h at the domain \mathcal{D} , denoted \mathcal{N} is developed first. There multiple ways to extend this to an approximation \tilde{h} on the full domain \mathcal{R} .

3.1 Function-Based Approximation: Method A

By using scaling relation (2.2.2) in the specific case of (2.2.3), we can reduce the dimension of the approximation problem from the full set \mathcal{R} to the domain \mathcal{D} .

Consider a neural network \mathcal{N} approximating h on the domain

$$\mathcal{D} = \{\mathbf{u} \in \mathcal{R} : u_0 = 1\}$$

$$\mathcal{N} : \mathcal{D} \longrightarrow \mathbb{R}$$

We make the definition

$$\tilde{h}_s := \mathcal{D} \longrightarrow \mathbb{R}$$

$$\tilde{h}(\mathbf{u}_N) := \mathcal{N}(\mathbf{u}_N)$$

and extend this to the full domain \mathcal{R} by the definition

$$\begin{aligned} \tilde{h} : \mathcal{R} &\longrightarrow \mathbb{R} \\ \tilde{h}(\mathbf{u}) &= u_0 \cdot \left(\tilde{h}_s\left(\frac{\mathbf{u}}{u_0}\right) + \log u_0 \right) \end{aligned}$$

In this case, we can calculate $\tilde{\alpha}$ from our approximation \tilde{h} :

$$\tilde{\alpha}(\mathbf{u}) := \nabla \tilde{h}(\mathbf{u})$$

For $\tilde{\alpha}$ we have the following identification:

$$\tilde{\alpha}^i(\mathbf{u}) = \begin{cases} \tilde{h}_s\left(\frac{\mathbf{u}}{u_0}\right) - \frac{\mathbf{u}_N \cdot \nabla_N \tilde{h}_s\left(\frac{\mathbf{u}}{u_0}\right)}{u_0} + \log u_0 + 1 & i = 0 \\ \partial_i \tilde{h}\left(\frac{\mathbf{u}}{u_0}\right) & i > 0 \end{cases}$$

3.2 Gradient-Based Approximation: Method B

Recall the following facts from the Lagrange-duality of the entropy-minimization framework:

$$\begin{aligned} h(\mathbf{u}) &= \int_{-1}^1 (\hat{\alpha}^t(\mathbf{u}) \mathbf{m} - 1) \exp(\hat{\alpha}^t(\mathbf{u}) \mathbf{m}) \\ \nabla h(\mathbf{u}) &= \hat{\alpha}(\mathbf{u}) \end{aligned}$$

$$\hat{u}_0(\alpha) = 1 \iff \alpha_0 = -\log \int_{-1}^1 \exp(\alpha_N^t \mathbf{m}_N) dt$$

Examining (2.2.1), it is clear that an approximation $\tilde{\alpha}_s$ to the map $\hat{\alpha}$ on the set \mathcal{D} , can be extended to \mathcal{R} via the log scaling relationship in (2.2.1). To obtain such an initial approximation $\tilde{\alpha}_s$, we use the gradient of a neural network \mathcal{N} approximating the function h at the realizable set.

First, assume that we have a neural network \mathcal{N} approximating h at the set \mathcal{D} to define $\tilde{\alpha}_s(\mathbf{u})$.

Definition. *Given a convex neural network \mathcal{N} approximating h at the set \mathcal{D} , we make the definition*

$$\tilde{\alpha}_s: \mathcal{D} \longrightarrow \mathbb{R}^{N+1}$$

$$\tilde{\alpha}_s^i(\mathbf{u}) = \begin{cases} -\log \int_{-1}^1 \exp(\nabla \mathcal{N} \cdot \mathbf{m}_N) dt & i = 0 \\ \nabla^i \mathcal{N}(\mathbf{u}) & i > 0 \end{cases} \quad (3.2.1)$$

Extend $\tilde{\alpha}_s$ to the whole domain \mathcal{R} according to (2.2.1) via the definition:

Definition. *($\tilde{\alpha}$) Define $\tilde{\alpha}: \mathcal{R} \longrightarrow \mathbb{R}^{N+1}$ via*

$$\tilde{\alpha} = \tilde{\alpha}_s\left(\frac{\mathbf{u}_N}{u_0}\right) + \begin{pmatrix} \log u_0 \\ \mathbf{0}_N \end{pmatrix}$$

Observe that this definition has the advantage that the moment-reconstruction map always reproduces the u_0 component exactly:

$$\tilde{\alpha}(\mathbf{u}) \rightarrow \hat{\mathbf{u}}(\tilde{\alpha}(\mathbf{u})) = \begin{pmatrix} u_0 \\ \hat{\mathbf{u}}_N(\tilde{\alpha}(\mathbf{u})) \end{pmatrix}$$

For our kinetic solution scheme, we seek to ensure that this definition of $\tilde{\alpha}$ be identifiable with the gradient of some convex function. For this, no global approximation to the entropy function \tilde{h} is needed, but one could additionally use this to define an approximation to the entropy function $\tilde{h}: \mathcal{R} \longrightarrow \mathbb{R}$ as

$$\tilde{h}(\mathbf{u}) = \int_{-1}^1 (\tilde{\alpha}^T \mathbf{m}(v) - 1) \exp(\tilde{\alpha}^T(\mathbf{u}) \mathbf{m}(v)) dv$$

3.3 Convexity of Method A

We want to know the requirements on \tilde{h}_s to ensure that \tilde{h} is a convex function, in the framework where:

$$\tilde{h}(\mathbf{u}) = u_0(\ln(u_0) + \tilde{h}_s\left(\frac{\mathbf{u}_N}{u_0}\right))$$

We make the definition

$$\xi(\mathbf{u}) := u_0 \cdot \tilde{h}_s\left(\frac{\mathbf{u}_N}{u_0}\right)$$

and note that, $\tilde{h}(\mathbf{u}) = u_0 \ln(u_0) + \xi(\mathbf{u})$. Since $u_0 \ln(u_0)$ is convex, it would suffice to require that ξ be convex in \mathcal{R} (since the sum of convex functions is again convex). In turn, to ensure convexity of ξ on \mathcal{R} , it suffices that \tilde{h}_s is convex and C^2 on its domain \mathcal{D} .

Fact. If \tilde{h}_s is convex on \mathcal{D} , then the map $\xi: \mathcal{R} \rightarrow \mathbb{R}$ defined by

$$\xi(\mathbf{u}) = u_0 \cdot \tilde{h}_s\left(\frac{\mathbf{u}_N}{u_0}\right)$$

is convex

Proof. As ξ is C^2 (since \tilde{h} is C^2), to prove that ξ is convex, we calculate the hessian and demonstrate that it is non-negative definite at all $\mathbf{u} \in \mathcal{R}$. Writing \mathbf{D}^2 and \mathbf{H} for the total second derivative matrices on \mathbb{R}^N and \mathbb{R}^{N+1} respectively, $\mathbf{H}(\xi)$ is the symmetric matrix with the following components

$$\mathbf{H}_{i,j}(\xi)(\mathbf{u}) = \begin{cases} \frac{1}{u_0^3} \langle \mathbf{u}_N^T, (\mathbf{D}^2 h_s)\left(\frac{\mathbf{u}}{u_0}\right) \cdot \mathbf{u}_N \rangle & i, j = 0 \\ -\frac{u_N}{u_0^2} (\mathbf{D}^2 h_s)_j\left(\frac{\mathbf{u}_N}{u_0}\right) & i = 0, j > 0 \\ \frac{1}{u_0} (\partial_{u_i, u_j} h_s)\left(\frac{\mathbf{u}_N}{u_0}\right) & i, j > 0 \end{cases}$$

To see that $\mathbf{H}(\xi)$ is positive definite on \mathcal{R} , first identify the following components of $\mathbf{H}(\xi)$:

$$a := \mathbf{H}_{0,0}(\xi)$$

$$b_i := \mathbf{H}_{0,i}(\xi)(\mathbf{u}), \quad i > 0$$

$$\mathbf{M}_{i,j} := \mathbf{H}_{i,j}(\xi), \quad i, j > 0$$

Since \tilde{h}_s is convex and C^2 , and $\frac{1}{u_0}$ is positive for all $\mathbf{u} \in \mathcal{R}$, the matrix $\mathbf{M} = (\mathbf{M}_{i,j})_{i,j>0}$ is non-negative definite. With this definition of \mathbf{M} observe that

$$a = \frac{\mathbf{u}_N^T}{u_0} \cdot \mathbf{M} \cdot \frac{\mathbf{u}_N}{u_0}$$

With these identifications available, applying $\mathbf{H}(\xi)(\mathbf{u})$ to a test vector $\mathbf{x} \in \mathbb{R}^{N+1}$ as a quadratic form gives

$$\mathbf{x}^T \cdot \mathbf{H}(\xi) \cdot \mathbf{x} = x_0^2 \cdot a + 2x_0 \sum_{i=1}^N x_i b_i + \mathbf{x}_N^T \cdot \mathbf{M} \cdot \mathbf{x}_N$$

Finally, replace a and b_i with their identities and let $\mathbf{v} = \frac{x_0}{u_0} \mathbf{u}_N - \mathbf{x}_N \in \mathbb{R}^N$ to see that

$$\begin{aligned} \mathbf{x}^T \cdot \mathbf{H}(\xi)(\mathbf{u}) \cdot \mathbf{x} &= \frac{x_0^2}{u_0^2} (\mathbf{u}_N^T \cdot \mathbf{M} \cdot \mathbf{u}_N) - 2 \frac{x_0}{u_0} (\mathbf{x}_N \cdot \mathbf{M} \cdot \mathbf{u}_N) + (\mathbf{x}_N \cdot \mathbf{M} \cdot \mathbf{x}_N) \\ &= \mathbf{v}^T \cdot \mathbf{M} \cdot \mathbf{v} \geq 0 \end{aligned}$$

since M is non-negative definite. As \mathbf{x} is arbitrary the proof is complete. \square

3.4 Convexity Question of Method B

We want to answer the following questions:

1. What are sufficient conditions on \mathcal{N} such that $\tilde{\alpha}(\mathbf{u})$ is equal to the gradient of a convex function?
2. What are sufficient conditions on \mathcal{N} such that \tilde{h} is a convex function?
3. What is the relation between $\nabla \tilde{h}$ and $\tilde{\alpha}$ (since they are distinct here)?

Answer To Question (1): We restrict to the case that \mathcal{N} is $C^2(\mathcal{D})$ so that by the Poincare Lemma for 1-forms, i.e. vector fields, we obtain a simple sufficient condition:

Fact. (*Convex Antiderivatives*) Let \mathbf{J} denote the total derivative ("jacobian") matrix for $N + 1$ dimensional vector field on \mathcal{R} , and \mathbf{D}^2 the total second-derivative ("hessian") matrix for C^2 functions on \mathcal{R} . If $\tilde{\alpha}: \mathcal{R} \rightarrow \mathbb{R}^{N+1}$ is a C^1 vector field with $\mathbf{J} = \mathbf{J}^T$ and

$$\mathbf{J}(\tilde{\alpha})(\mathbf{u}) \geq 0 \quad \forall \mathbf{u} \in \mathcal{R} \text{ (in the sense of matrices)}$$

then \exists convex $f \in C^2(\mathcal{R}, \mathbb{R})$ with $\nabla f = \tilde{\alpha}$.

This fact says that $\tilde{\alpha}$ is the gradient of a convex function provided that \mathbf{J} is symmetric, and $\mathbf{J}(\tilde{\alpha}) \geq 0$, so we seek sufficient conditions on \mathcal{N} such that $\mathbf{J}(\tilde{\alpha}) \geq 0$ and $\mathbf{J} = \mathbf{J}^T$ by calculating $\mathbf{J}(\tilde{\alpha})$ explicitly (again, where $\tilde{\alpha}$ is defined via a neural network exactly as in (3.2.1). To simplify this calculation we define the **moment-reconstruction** induced by $\tilde{\alpha}$ as

$$\hat{\mathbf{u}}(\tilde{\alpha}(\mathbf{u}))$$

and note that this vector quantity satisfies the equality

$$\frac{\hat{\mathbf{u}}_N(\tilde{\alpha}(\mathbf{u}))}{\hat{u}_0(\tilde{\alpha}(\mathbf{u}))} = \frac{\int_{-1}^1 \exp(\nabla \mathcal{N}(\frac{\mathbf{u}_N}{u_0}) \cdot \mathbf{m}_N) \mathbf{m}_N(v) dv}{\int_{-1}^1 \exp(\nabla \mathcal{N}(\frac{\mathbf{u}_N}{u_0}) \mathbf{m}_N) dv}$$

Using these identities, and writing $\hat{\mathbf{u}}$ in lieu of $\hat{\mathbf{u}}(\tilde{\alpha}(\mathbf{u}))$, it follows that $\mathbf{J}(\tilde{\alpha})$ is the the matrix with the following components:

$$\mathbf{J}_{i,j}(\tilde{\alpha})(\mathbf{u}) = \begin{cases} \frac{1}{u_0} + \frac{1}{u_0} (\frac{\mathbf{u}_N}{u_0} \cdot \mathbf{D}^2 \mathcal{N}(\frac{\mathbf{u}}{u_0}) \cdot \frac{\hat{\mathbf{u}}_N}{\hat{u}_0}) & i, j = 0 \\ \frac{-1}{u_0} \cdot \mathbf{D}_i^2(\mathcal{N})(\frac{\mathbf{u}_N}{u_0}) \cdot \frac{\hat{\mathbf{u}}_N}{\hat{u}_0} & i = 0, j > 0 \\ \frac{-1}{u_0} \cdot \mathbf{D}_i^2(\mathcal{N})(\frac{\mathbf{u}_N}{u_0}) \cdot \frac{\mathbf{u}_N}{u_0} & i > 0, j = 0 \\ \frac{1}{u_0} \mathbf{D}_{i,j}^2 \mathcal{N}(\frac{\mathbf{u}}{u_0}) & i, j > 0 \end{cases}$$

Note the asymmetry in the matrix $\mathbf{J}(\tilde{\alpha})$ occuring in the first column, and first row, due to the $\hat{\mathbf{u}}_N/\hat{u}_0$ vector. In the case where the reconstruction $\hat{\mathbf{u}}$ is "exact", then we would have that $\mathbf{J}(\tilde{\alpha})$ is not symmetric and therefore $\tilde{\alpha}$ is not the gradient of any function, let alone a convex one.